

Kubernetes Deployment Guide for Archivist

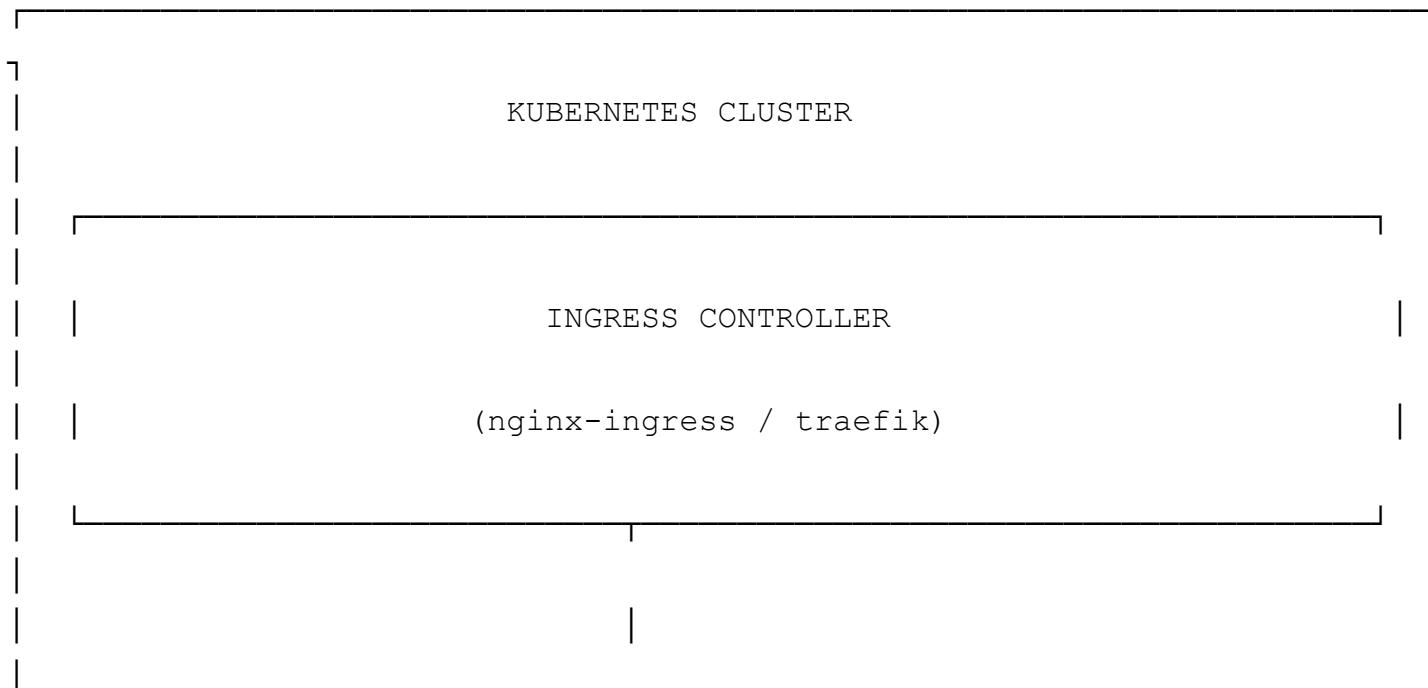
This guide explains how to deploy the Archivist application on Kubernetes for production-grade scalability, high availability, and orchestration.

Table of Contents

1. [Architecture Overview](#)
 2. [Why Kubernetes for Archivist](#)
 3. [Prerequisites](#)
 4. [Kubernetes Components](#)
 5. [Deployment Strategy](#)
 6. [Service Definitions](#)
 7. [Configuration Management](#)
 8. [Scaling Strategies](#)
 9. [Monitoring and Observability](#)
 10. [Production Considerations](#)
-

Architecture Overview

Archivist's microservices architecture maps naturally to Kubernetes:



APPLICATION NAMESPACE

```
RAG API      | Graph       | Archivist Worker  
Deployment   | Service     | Jobs / CronJobs  
(FastAPI)   | Deployment | (PDF Processing)  
Replicas: 3 | Replicas: 2 |
```

DATA NAMESPACE

```
Neo4j        | Qdrant      | Redis       | Kafka  
StatefulSet | StatefulSet | StatefulSet | (Strimzi)  
(Graph DB) | (Vector DB) | (Cache)    |
```

PERSISTENT STORAGE

StorageClass: SSD (gp3/premium-lrs) for databases

PersistentVolumeClaims for: Neo4j, Qdrant, Redis, PDF storage

Why Kubernetes for Archivist

Current Challenges (Docker Compose)

| Challenge | Impact |
|------------------------|---|
| Single-node deployment | No fault tolerance |
| Manual scaling | Cannot handle traffic spikes |
| No self-healing | Service failures require manual restart |
| Resource contention | All services share same host resources |
| No rolling updates | Downtime during deployments |

Kubernetes Benefits

| Benefit | Archivist Use Case |
|-----------------------------------|---|
| Horizontal Pod Autoscaling | Scale RAG API pods during high query load |
| StatefulSets | Reliable Neo4j, Qdrant, Redis deployments with persistent storage |
| Jobs/CronJobs | Batch PDF processing, scheduled index rebuilds |
| Service Discovery | Automatic DNS for inter-service communication |
| ConfigMaps/Secrets | Centralized config, secure API key management |
| Health Checks | Auto-restart failed containers |
| Resource Quotas | Prevent runaway Gemini API calls |
| Rolling Updates | Zero-downtime deployments |
| Multi-zone HA | Survive data center failures |

Prerequisites

Required Tools

```
# kubectl - Kubernetes CLI  
curl -LO "https://dl.k8s.io/release/$(curl -L -s https://dl.k8s.io/releas  
  
# helm - Package manager for Kubernetes  
curl https://raw.githubusercontent.com/helm/helm/main/scripts/get-helm-3  
  
# Optional: k9s for cluster management  
brew install derailed/k9s/k9s
```

Cluster Options

| Option | Best For | Notes |
|-------------------|-------------------|---------------------------------------|
| Minikube | Local development | Single-node, limited resources |
| kind | CI/CD testing | Docker-based, ephemeral |
| EKS(AWS) | Production | Managed, integrates with AWS services |
| GKE(GCP) | Production | Best for Gemini API (same network) |
| AKS(Azure) | Production | Managed, enterprise features |

Kubernetes Components

Namespace Organization

```
# namespaces.yaml  
apiVersion: v1  
kind: Namespace  
metadata:  
  name: archivist  
  labels:  
    app: archivist  
    environment: production  
---
```

```
apiVersion: v1
kind: Namespace
metadata:
  name: archivist-data
  labels:
    app: archivist
    tier: data
```

1. RAG API Service (Deployment)

```
# rag-api-deployment.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: rag-api
  namespace: archivist
  labels:
    app: rag-api
    tier: api
spec:
  replicas: 3
  selector:
    matchLabels:
      app: rag-api
  template:
    metadata:
      labels:
        app: rag-api
    spec:
      containers:
        - name: rag-api
          image: archivist/rag-api:latest
          ports:
            - containerPort: 8000
          env:
            - name: GEMINI_API_KEY
              valueFrom:
                secretKeyRef:
                  name: archivist-secrets
                  key: gemini-api-key
            - name: QDRANT_HOST
              value: "qdrant.archivist-data.svc.cluster.local"
```

```

- name: QDRANT_PORT
  value: "6333"
- name: REDIS_HOST
  value: "redis.archivist-data.svc.cluster.local"
- name: NEO4J_URI
  value: "bolt://neo4j.archivist-data.svc.cluster.local:7687"
resources:
  requests:
    memory: "512Mi"
    cpu: "250m"
  limits:
    memory: "2Gi"
    cpu: "1000m"
livenessProbe:
  httpGet:
    path: /health
    port: 8000
  initialDelaySeconds: 30
  periodSeconds: 10
readinessProbe:
  httpGet:
    path: /health
    port: 8000
  initialDelaySeconds: 5
  periodSeconds: 5
---
apiVersion: v1
kind: Service
metadata:
  name: rag-api
  namespace: archivist
spec:
  selector:
    app: rag-api
  ports:
  - port: 8000
    targetPort: 8000
  type: ClusterIP

```

2. Graph Service (Deployment)

```
# graph-service-deployment.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: graph-service
  namespace: archivist
  labels:
    app: graph-service
spec:
  replicas: 2
  selector:
    matchLabels:
      app: graph-service
  template:
    metadata:
      labels:
        app: graph-service
    spec:
      containers:
        - name: graph-service
          image: archivist/graph-service:latest
          ports:
            - containerPort: 8081
          env:
            - name: GEMINI_API_KEY
              valueFrom:
                secretKeyRef:
                  name: archivist-secrets
                  key: gemini-api-key
            - name: NEO4J_URI
              value: "bolt://neo4j.archivist-data.svc.cluster.local:7687"
            - name: NEO4J_USER
              value: "neo4j"
            - name: NEO4J_PASSWORD
              valueFrom:
                secretKeyRef:
                  name: archivist-secrets
                  key: neo4j-password
            - name: KAFKA_BOOTSTRAP_SERVERS
              value: "kafka-cluster-kafka-bootstrap.archivist-data.svc.cluster.local:9092"
      resources:
        requests:
          memory: "256Mi"
          cpu: "100m"
```

```

limits:
  memory: "1Gi"
  cpu: "500m"
livenessProbe:
  httpGet:
    path: /health
    port: 8081
  initialDelaySeconds: 30
  periodSeconds: 10

```

3. Neo4j (StatefulSet)

```

# neo4j-statefulset.yaml
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: neo4j
  namespace: archivist-data
spec:
  serviceName: neo4j
  replicas: 1  # Use Neo4j Cluster for HA (3+ replicas)
  selector:
    matchLabels:
      app: neo4j
  template:
    metadata:
      labels:
        app: neo4j
    spec:
      containers:
        - name: neo4j
          image: neo4j:5.13-community
          ports:
            - containerPort: 7474
              name: http
            - containerPort: 7687
              name: bolt
      env:
        - name: NEO4J_AUTH
          valueFrom:
            secretKeyRef:
              name: neo4j-credentials

```

```

        key: auth
      - name: NEO4J_PLUGINS
        value: '["graph-data-science", "apoc"]'
      - name: NEO4J_dbms_memory_heap_max_size
        value: "2G"
    volumeMounts:
      - name: neo4j-data
        mountPath: /data
      - name: neo4j-logs
        mountPath: /logs
  resources:
    requests:
      memory: "2Gi"
      cpu: "500m"
    limits:
      memory: "4Gi"
      cpu: "2000m"
  volumeClaimTemplates:
    - metadata:
        name: neo4j-data
      spec:
        accessModes: ["ReadWriteOnce"]
        storageClassName: ssd
        resources:
          requests:
            storage: 50Gi
    - metadata:
        name: neo4j-logs
      spec:
        accessModes: ["ReadWriteOnce"]
        storageClassName: standard
        resources:
          requests:
            storage: 10Gi
  ---
apiVersion: v1
kind: Service
metadata:
  name: neo4j
  namespace: archivist-data
spec:
  selector:
    app: neo4j
  ports:

```

```
- port: 7474
  targetPort: 7474
  name: http
- port: 7687
  targetPort: 7687
  name: bolt
type: ClusterIP
```

4. Qdrant (StatefulSet)

```
# qdrant-statefulset.yaml
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: qdrant
  namespace: archivist-data
spec:
  serviceName: qdrant
  replicas: 1  # Can scale to 3+ for distributed mode
  selector:
    matchLabels:
      app: qdrant
  template:
    metadata:
      labels:
        app: qdrant
    spec:
      containers:
        - name: qdrant
          image: qdrant/qdrant:v1.7.4
          ports:
            - containerPort: 6333
              name: http
            - containerPort: 6334
              name: grpc
          volumeMounts:
            - name: qdrant-storage
              mountPath: /qdrant/storage
      resources:
        requests:
          memory: "1Gi"
          cpu: "250m"
```

```

limits:
  memory: "4Gi"
  cpu: "1000m"
livenessProbe:
  httpGet:
    path: /healthz
    port: 6333
  initialDelaySeconds: 30
  periodSeconds: 10
volumeClaimTemplates:
- metadata:
  name: qdrant-storage
spec:
  accessModes: ["ReadWriteOnce"]
  storageClassName: ssd
  resources:
    requests:
      storage: 100Gi
---
apiVersion: v1
kind: Service
metadata:
  name: qdrant
  namespace: archivist-data
spec:
  selector:
    app: qdrant
  ports:
  - port: 6333
    targetPort: 6333
    name: http
  - port: 6334
    targetPort: 6334
    name: grpc

```

5. Redis (StatefulSet)

```

# redis-statefulset.yaml
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: redis

```

```
namespace: archivist-data
spec:
  serviceName: redis
  replicas: 1
  selector:
    matchLabels:
      app: redis
  template:
    metadata:
      labels:
        app: redis
    spec:
      containers:
        - name: redis
          image: redis:7.2-alpine
          ports:
            - containerPort: 6379
          command:
            - redis-server
            - --appendonly
            - "yes"
            - --maxmemory
            - "512mb"
            - --maxmemory-policy
            - "allkeys-lru"
          volumeMounts:
            - name: redis-data
              mountPath: /data
      resources:
        requests:
          memory: "256Mi"
          cpu: "100m"
        limits:
          memory: "1Gi"
          cpu: "500m"
      livenessProbe:
        exec:
          command:
            - redis-cli
            - ping
        initialDelaySeconds: 30
        periodSeconds: 10
  volumeClaimTemplates:
    - metadata:
```

```

name: redis-data
spec:
  accessModes: ["ReadWriteOnce"]
  storageClassName: ssd
  resources:
    requests:
      storage: 10Gi
---
apiVersion: v1
kind: Service
metadata:
  name: redis
  namespace: archivist-data
spec:
  selector:
    app: redis
  ports:
    - port: 6379
      targetPort: 6379

```

6. PDF Processing Job

```

# pdf-processor-job.yaml
apiVersion: batch/v1
kind: Job
metadata:
  name: pdf-processor-batch
  namespace: archivist
spec:
  parallelism: 4 # Process 4 PDFs concurrently
  completions: 10 # Total PDFs to process
  backoffLimit: 3
  template:
    spec:
      containers:
        - name: processor
          image: archivist/archivist:latest
          command: ["./archivist", "process", "--batch"]
          env:
            - name: GEMINI_API_KEY
              valueFrom:
                secretKeyRef:

```

```

        name: archivist-secrets
        key: gemini-api-key
volumeMounts:
- name: pdf-storage
  mountPath: /app/lib
- name: output-storage
  mountPath: /app/tex_files
resources:
  requests:
    memory: "1Gi"
    cpu: "500m"
  limits:
    memory: "4Gi"
    cpu: "2000m"
volumes:
- name: pdf-storage
  persistentVolumeClaim:
    claimName: pdf-input-pvc
- name: output-storage
  persistentVolumeClaim:
    claimName: latex-output-pvc
restartPolicy: OnFailure

```

7. Index Rebuild CronJob

```

# index-rebuild-cronjob.yaml
apiVersion: batch/v1
kind: CronJob
metadata:
  name: index-rebuild
  namespace: archivist
spec:
  schedule: "0 2 * * *"  # Daily at 2 AM
  jobTemplate:
    spec:
      template:
        spec:
          containers:
            - name: indexer
              image: archivist/rag-api:latest
              command: ["python", "-m", "indexer", "--rebuild"]
          env:

```

```

- name: GEMINI_API_KEY
  valueFrom:
    secretKeyRef:
      name: archivist-secrets
      key: gemini-api-key
- name: QDRANT_HOST
  value: "qdrant.archivist-data.svc.cluster.local"
resources:
  requests:
    memory: "2Gi"
    cpu: "500m"
  limits:
    memory: "8Gi"
    cpu: "2000m"
restartPolicy: OnFailure

```

Configuration Management

Secrets

```

# secrets.yaml
apiVersion: v1
kind: Secret
metadata:
  name: archivist-secrets
  namespace: archivist
type: Opaque
stringData:
  gemini-api-key: "your-gemini-api-key"
  neo4j-password: "your-neo4j-password"
---
apiVersion: v1
kind: Secret
metadata:
  name: neo4j-credentials
  namespace: archivist-data
type: Opaque
stringData:
  auth: "neo4j/your-neo4j-password"

```

ConfigMap

```
# configmap.yaml
apiVersion: v1
kind: ConfigMap
metadata:
  name: archivist-config
  namespace: archivist
data:
  config.yaml: |
    processing:
      max_workers: 8
      batch_size: 10
      timeout_per_paper: 600s

  gemini:
    model: "models/gemini-2.0-flash-exp"
    temperature: 0.3
    agentic_workflow: true

  graph:
    async_building: true
    max_graph_workers: 2
    search:
      vector_weight: 0.5
      graph_weight: 0.3
      keyword_weight: 0.2

  qdrant:
    collection_name: "archivist_papers"
    vector_size: 768
    distance: "Cosine"

  cache:
    ttl_days: 30
```

Scaling Strategies

Horizontal Pod Autoscaler (HPA)

```

# hpa.yaml
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: rag-api-hpa
  namespace: archivist
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: rag-api
  minReplicas: 2
  maxReplicas: 10
  metrics:
    - type: Resource
      resource:
        name: cpu
        target:
          type: Utilization
          averageUtilization: 70
    - type: Resource
      resource:
        name: memory
        target:
          type: Utilization
          averageUtilization: 80
  behavior:
    scaleDown:
      stabilizationWindowSeconds: 300
      policies:
        - type: Percent
          value: 10
          periodSeconds: 60
    scaleUp:
      stabilizationWindowSeconds: 0
      policies:
        - type: Percent
          value: 100
          periodSeconds: 15

```

Vertical Pod Autoscaler (VPA)

```
# vpa.yaml
apiVersion: autoscaling.k8s.io/v1
kind: VerticalPodAutoscaler
metadata:
  name: rag-api-vpa
  namespace: archivist
spec:
  targetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: rag-api
  updatePolicy:
    updateMode: "Auto"
  resourcePolicy:
    containerPolicies:
      - containerName: rag-api
        minAllowed:
          cpu: 100m
          memory: 256Mi
        maxAllowed:
          cpu: 4
          memory: 8Gi
```

Monitoring and Observability

Prometheus ServiceMonitor

```
# servicemonitor.yaml
apiVersion: monitoring.coreos.com/v1
kind: ServiceMonitor
metadata:
  name: archivist-metrics
  namespace: archivist
spec:
  selector:
    matchLabels:
      app: rag-api
  endpoints:
    - port: http
```

```
path: /metrics
interval: 30s
```

Key Metrics to Monitor

| Service | Metric | Alert Threshold |
|---------|----------------------|-----------------|
| RAG API | Response latency p99 | > 2s |
| RAG API | Error rate | > 5% |
| Qdrant | Memory usage | > 80% |
| Neo4j | Active connections | > 100 |
| Redis | Memory fragmentation | > 1.5 |
| Kafka | Consumer lag | > 1000 |
| All | Pod restarts | > 3/hour |

Production Considerations

High Availability Setup

Recommended Production Configuration:

- RAG API: 3+ replicas across availability zones
- Graph Service: 2+ replicas
- Neo4j: 3-node cluster (Enterprise)
- Qdrant: 3-node distributed cluster
- Redis: 3-node Sentinel or Redis Cluster
- Kafka: 3+ brokers via Strimzi operator

Resource Allocation Guide

| Component | CPU Request | CPU Limit | Memory Request | Memory Limit | Storage |
|---------------|-------------|-----------|----------------|--------------|-----------|
| RAG API | 250m | 1000m | 512Mi | 2Gi | - |
| Graph Service | 100m | 500m | 256Mi | 1Gi | - |
| Neo4j | 500m | 2000m | 2Gi | 4Gi | 50Gi SSD |
| Qdrant | 250m | 1000m | 1Gi | 4Gi | 100Gi SSD |

| Component | CPU Request | CPU Limit | Memory Request | Memory Limit | Storage |
|---------------|-------------|-----------|----------------|--------------|----------|
| Redis | 100m | 500m | 256Mi | 1Gi | 10Gi SSD |
| PDF Processor | 500m | 2000m | 1Gi | 4Gi | - |

Network Policies

```
# network-policy.yaml
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: rag-api-policy
  namespace: archivist
spec:
  podSelector:
    matchLabels:
      app: rag-api
  policyTypes:
  - Ingress
  - Egress
  ingress:
  - from:
    - namespaceSelector:
        matchLabels:
          name: ingress-nginx
    ports:
    - port: 8000
  egress:
  - to:
    - namespaceSelector:
        matchLabels:
          app: archivist
  - to:
    - namespaceSelector:
        matchLabels:
          tier: data
  # Allow external access to Gemini API
  - to:
    - ipBlock:
        cidr: 0.0.0.0/0
  ports:
  - port: 443
```

Ingress Configuration

```
# ingress.yaml
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: archivist-ingress
  namespace: archivist
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /
    cert-manager.io/cluster-issuer: letsencrypt-prod
spec:
  ingressClassName: nginx
  tls:
  - hosts:
    - api.archivist.example.com
    secretName: archivist-tls
  rules:
  - host: api.archivist.example.com
    http:
      paths:
      - path: /rag
        pathType: Prefix
        backend:
          service:
            name: rag-api
            port:
              number: 8000
      - path: /graph
        pathType: Prefix
        backend:
          service:
            name: graph-service
            port:
              number: 8081
```

Deployment Commands

Initial Setup

```

# Create namespaces
kubectl apply -f namespaces.yaml

# Deploy secrets (use sealed-secrets or external-secrets in production)
kubectl apply -f secrets.yaml

# Deploy ConfigMaps
kubectl apply -f configmap.yaml

# Deploy data layer
kubectl apply -f neo4j-statefulset.yaml
kubectl apply -f qdrant-statefulset.yaml
kubectl apply -f redis-statefulset.yaml

# Wait for data services to be ready
kubectl wait --for=condition=ready pod -l app=neo4j -n archivist-data --t
kubectl wait --for=condition=ready pod -l app=qdrant -n archivist-data --
kubectl wait --for=condition=ready pod -l app=redis -n archivist-data --t

# Deploy application layer
kubectl apply -f rag-api-deployment.yaml
kubectl apply -f graph-service-deployment.yaml

# Deploy autoscaling
kubectl apply -f hpa.yaml

# Deploy ingress
kubectl apply -f ingress.yaml

```

Useful Commands

```

# Check pod status
kubectl get pods -n archivist
kubectl get pods -n archivist-data

# View logs
kubectl logs -f deployment/rag-api -n archivist

# Scale manually
kubectl scale deployment rag-api --replicas=5 -n archivist

# Run one-off PDF processing job

```

```
kubectl create job --from=cronjob/pdf-processor pdf-batch-$($date +%s) -n archivist

# Port-forward for local access
kubectl port-forward svc/rag-api 8000:8000 -n archivist
kubectl port-forward svc/neo4j 7474:7474 7687:7687 -n archivist-data

# Check HPA status
kubectl get hpa -n archivist

# View resource usage
kubectl top pods -n archivist
```

Helm Chart Structure (Recommended)

For production deployments, package the manifests as a Helm chart:

```
archivist-chart/
├── Chart.yaml
├── values.yaml
├── values-production.yaml
└── templates/
    ├── _helpers.tpl
    ├── namespaces.yaml
    ├── secrets.yaml
    ├── configmap.yaml
    ├── rag-api/
    │   ├── deployment.yaml
    │   ├── service.yaml
    │   └── hpa.yaml
    ├── graph-service/
    │   ├── deployment.yaml
    │   └── service.yaml
    ├── neo4j/
    │   ├── statefulset.yaml
    │   └── service.yaml
    ├── qdrant/
    │   ├── statefulset.yaml
    │   └── service.yaml
    ├── redis/
    │   ├── statefulset.yaml
    │   └── service.yaml
```

```
|   └── jobs/
|       ├── pdf-processor.yaml
|       └── index-rebuild.yaml
|   └── ingress.yaml
└── networkpolicy.yaml
└── README.md
```

Install with:

```
helm install archivist ./archivist-chart -f values-production.yaml
```

Migration from Docker Compose

1. **Build container images** and push to a registry (Docker Hub, ECR, GCR)
 2. **Export data** from local volumes to cloud storage or PVCs
 3. **Test in staging** cluster before production
 4. **Use blue-green deployment** for zero-downtime migration
 5. **Monitor closely** during initial production rollout
-

Cost Optimization Tips

1. Use **Spot/Preemptible instances** for PDF processing jobs
 2. Enable **cluster autoscaler** to scale down during off-peak hours
 3. Use **resource quotas** to prevent runaway costs
 4. Consider **serverless** options (Knative, Cloud Run) for bursty workloads
 5. Use **managed services** (Cloud Memorystore, Neo4j Aura) to reduce ops overhead
-

Next Steps

1. Set up CI/CD pipeline for automated deployments
2. Configure backup strategies for databases
3. Implement disaster recovery procedures
4. Set up log aggregation (ELK/Loki)
5. Configure alerting (PagerDuty/Slack integration)